

MAX 2007

CONNECT. DISCOVER. INSPIRE.

Alistair McLeod

Practice Leader

Adobe Consulting



Building RIAs with Cairngorm and LiveCycle Data Services



- Alistair McLeod
 - Adobe Consulting
 - Based in Edinburgh, Scotland
 - Practice Leader for Technical Practice in EMEA
- Background
 - J2EE Enterprise Software Development
 - Agile Software Development
 - Rich Internet Application Development
 - Co-founder of iteration::two
 - Acquired by Macromedia in 2005
 - Co-author “Developing Rich Clients with Macromedia Flex”
 - Core contributor to Cairngorm and FlexUnit Open Source Projects

What we'll cover in the next 60 minutes



- MFG.com project
 - Briefly set the context
- Cairngorm Enterprise
 - Share our current thinking
- Data Management Services
 - Best practices
 - Hints and tips

Who are you?



- Have you used Cairngorm?
 - Have you built your own framework or adapted Cairngorm to fit your own requirements?
- Have you used Data Management services?
 - Have you found the experiences positive?
- Do you understand Scottish?

What are my assumptions?



- ✓ You are familiar with the basics of Cairngorm
 - ✓ And the design patterns used
- ✓ You have used or understand the concepts of Data Management Services

What are the objectives of this session?



- ✔ Share our experiences of building large applications with Cairngorm
 - ✔ Particularly using Data Management Services
- ✔ Provide these learning experiences as best practices
- ✔ Base it on a real-world project (MFG.com)
- ✔ Look at how this experience may influence Cairngorm Enterprise
- ✔ Provide ideas on how you might approach your own projects
- ✔ Answer your questions
- ✘ We won't be providing any code

MFG.com project



- Global online manufacturing marketplace
 - Matches buyers and suppliers for the manufacturing of machined parts
- HQ in Atlanta, Georgia, USA
- Offices around the world, namely Geneva and Shanghai
 - Geneva office gained through acquisition of main competitor
- President, CEO and Founder is Mitch Free
 - Engaged with Adobe Consulting
- Biggest investor is Jeff Bezos (Amazon.com President, CEO and Founder)
- MFG Won Max award in Chicago
 - Communication and Collaboration

What is the MFG.com terminology?



- Buyer
 - Wants something manufactured
- Supplier
 - Manufactures things
- RFQ (Request for Quote)
 - Created by buyer
 - Made up of multiple parts
 - Contains part drawings and specifications
- Parts
 - Standard part (out of a brochure, eg, bolt, screw)
 - Custom part (specific to buyer, eg, phone cover, phone screen,)
- Quote
- Award

*A 10 point footnote can go here, if necessary

What are the goals of the project?



- MFG want to build a global platform to:
 - Consolidate US, European and Asian markets
 - Allow users to work in their native language, currency and date time formats
 - Accelerate market growth
- Utilize RIA to allow the delivery of complex business functionality to users
 - Users can be huge multinational organisations or companies with less than 5 employees
 - Allow user to achieve their goals more quickly
 - While retaining ease of use
- Reduce time to complete complex and repetitive tasks
- A platform on which to realize MFG's technology vision
 - Collaboration between buyers and suppliers
 - Interested in Pacifica (VOIP) and Cocomo (Collaboration components)

- Flex 2.0.1 hotfix 2
- LiveCycle Data Services (LC DS) 2.5
- JBoss Application Server 4.0.4
- Oracle database

Application demo



What are the technical challenges behind the app?



- There is a lot of data
 - You only saw a small part of the application
 - And its still growing
- There are many views on the data, which require a lot of associations
 - Complex business model
 - Different views on the same data
 - eg, Data can be viewed from an RFQ perspective, a part perspective or a quote perspective
 - Relevant when using Data Management Services
 - Associations defined in configuration files
- It soon becomes very complex
 - Can't use a single domain model (because of different associations and views onto the model)
 - We use presentation models rather than domain models (for views)
- There is a lot of functionality in the application

*A 10 point footnote can go here, if necessary

- Total number of users: ~250K
- Concurrent users: 250+
- Average number of open RFQs: 1500
- Average number of quotes per RFQ: 4 -6
- Total classes: ~1500
- Total Lines of Code : ~100K
- Number of data management services: ~65
 - Relatively large number due to number of different views on same data
 - Results in more data management services
- Still under development and growing in size!

Cairngorm Enterprise



- Cairngorm Enterprise already exists
 - Created in response to changes in Flex SDK Hotfix2/LiveCycle Data Services (LCDS) 2.5.x
 - Currently version 2.2.1
 - On Adobe Labs
- Created to remove Cairngorm dependency on LCDS
 - Consumer and Producer classes moved into LCDS
 - Can build applications upon standard Flex SDK
 - No dependency on LCDS
- Existing Cairngorm Enterprise supports services provided by LCDS
 - Provides EnterpriseServiceLocator for Consumer and Producer
 - Added no new functionality

Why consider changing Cairngorm Enterprise?



- Large-scale applications introduce new challenges
 - MFG.com project is **big**
 - Lots of events, commands, large controller
- Cairngorm today is very applicable to a Service Oriented Architecture (SOA)
- Data Management Services gives us a Data Oriented Architecture
 - How well does Cairngorm apply to a Data Oriented Architecture?

What are our goals of this new thinking?



- Recognize and exploit the strengths of the Flex programming model
 - eg, MXML and Binding
 - Current patterns in Cairngorm don't necessarily take that into account
- Don't want to force patterns on to Flex if they are not appropriate
 - But we must recognize peoples familiarity and trust of well known patterns
 - Provide a common vocabulary for development teams
- Make it easier to develop and maintain larger code bases
 - Make code simple and easy to understand
- Better support for LCDS-based applications
 - In particular, those utilising Data Management Services
 - Cairngorm built around SOA patterns, how does it apply to a Data Oriented Architecture?
- Based on our own experience from real world projects
- Listening to the community
 - Open up discussions

*A 10 point footnote can go here, if necessary

- Still very much work in progress
- The Adobe Consulting way is to “release once proven”
 - Identify recurring solutions from our own engagements
- We are not advocating new Cairngorm best-practices at this time
- Provoke discussion around Cairngorm Enterprise
- Influence how you approach your projects
- We want to share our ideas
- Solicit your feedback and your own ideas
- This is a sneak!!!
 - Don't expect a new release of Cairngorm based on these ideas anytime soon

- Using Cairngorm with large scale applications
- Model-View architecture
 - How well does MVC architecture apply to Flex development
 - What benefits is the controller bringing?
- Data Management Services
 - How does Cairngorm and Flex fit within a Data Oriented Architecture
- Security
- Some other bits and pieces
 - We'll show some best practices

- A new package structure
 - `com.adobe.cairngorm.enterprise.control`
 - Controller, command and event classes
 - `com.adobe.cairngorm.enterprise.model`
 - Model classes
 - `com.adobe.cairngorm.enterprise.service`
 - Service and responder classes

- ✓ More descriptive and concise top-level organization

- Promote a good OO model
 - Describes the state and behavior of the application
 - Don't assume all logic has to go into a "Cairngorm Class"
- The view binds directly onto the model (as is already the case with Cairngorm)
 - At the moment the ModelLocator is synonymous to the HTTP session of the web world
- User gestures executed directly onto the model
- Move the Controller "behind" the model – see next slide

- ✓ More concise code / less verbosity
- ✓ Application structure will be cleaner and easier to maintain and extend
- ✓ Encourage developers to move logic out of the views and in to the model
- ✓ Make unit testing and test driven development easier

- ✗ This breaks MVC
- ✗ User gestures no longer mapped to Cairngorm events
- ✗ No final decision made yet

“Dependency injection”



- The model is injected into command and the responder
 - More explicit separation of command and responder
- The ServiceLocator is injected into the delegate.
- The basic principle is the model makes calls to the service, via the controller
 - Dispatches an event, which is picked up by the controller and passed onto a command and delegate
 - Model can then be updated with the result in the responder

- ✓ Gives the usual documented benefits of Dependency Injection. Also...
- ✓ Easier to test (inject a mock ServiceLocator into delegate)
- ✓ Avoids excessive use of ModelLocator
- ✓ Avoids lengthy dot notations

- ✗ What benefit is the controller bringing?
- ✗ Replace with Delegate Factory?
 - ✗ Would also remove events and commands

- This is where we would like to be
 - Turning off security should be a conscious decision by the developer
 - Make security inherent within Cairngorm Enterprise applications
 - Build on security model provided by LC DS and the underlying application server
 - Hooks into setCredentials()
-
- ✓ Bring repeatability and re-usability to login and logout
 - ✓ Saves development effort
 - ✓ Consistent approach to security
 - ✓ Quicker for developers to get up to speed

- Support for dealing with `ItemPendingError`
 - Thrown when data not yet retrieved from server – also causes request for that data
 - Results in lots of try/catch blocks in code – reduces code clarity
 - Created `ItemPendingResponder` to help, but not ideal
- Handling data conflicts
 - Only update/update conflicts are handled by data management services
 - Provide a consistent way to handle data errors
 - User experience may have “moved on” – how do we present errors to the user?
- Allow us to keep specific data on logout
 - Eg, Localisation data - don't want to fetch every time someone new logs in

✓ We want to help make developers successful with data management services and a data oriented architecture

- ✗ Need a better way of handling fault codes
 - ✗ if-then-else is clumsy
- ✗ Opportunity to provide better support around item pending errors

- Logging...
- Modules...
- AIR-specific Cairngorm plugins
 - eg, Implementation of DAO pattern for SQLite
- Any more?

Data Management Services Best Practices



Getters shouldn't update state



- You will see unexpected results if the class is introspected
 - eg, if logging enabled or during de-serialization
- Can cause you to get unexpected results when debugging
- This is a general principle, not specific to Data Management Services
 - But can be difficult to track down when using Data Management Services
 - A lot of code is auto-generated.

- A remote collection is, to the client code, an `ICollectionView`
 - i.e. the client code doesn't know if it is a local collection or remote collection.
- Consider the implications of using `sort` or `filterFunction` on a paged collection
 - Paged collections means that only some items reside on the client
 - Other items are fetched on-demand
 - But, sorting and filtering is undertaken on the client and needs all the items in the collection
- This removes the benefits paging
 - All data is retrieved on the client before sorting or filtering
 - May affect the user experience – user may have to wait for a response
- Let the server (and database) do its job
 - Use `DataService.fill()` to retrieve the data on a sort or filter
 - Pass sort and/or filter parameters

- One of benefits of a remote collection is paging
- You need to think about `ItemPendingError`
 - Thrown when accessing an object that doesn't yet exist on client
 - Have to catch error and add responder
 - Re-execute logic when responder is called
 - Leads to more complicated code
- Create separate responder classes to handle the results

```
try
{
    for ( var i : uint = 0; i < employees.length; i++ )
    {
        var employee : Employee = Employee( employees.getItemAt( i ) );
        employee.fire();
    }
}
catch( error : ItemPendingError )
{
    error.addResponder( new ItemPendingResponder( Employee, "fire" ) );
}
```

- An MXML binding expression provides an elegant way to handle an `ItemPendingError`
 - Catches the `ItemPendingError`
 - Adds responder
 - Fires change event when responder is called
 - Your view is updated automatically
- This removes the need to write exception handling code
 - No need to create responder
- Simplifies your development
 - Less code to write
 - Improves code clarity
- Especially useful in custom item renderers

```
<mx:Label text= "{ employees.getItem( 1 ).address.postcode }"/>
```


Set autoCommit() to 'false'



- By default autoCommit is set to true on DataService
- Every change to a managed object results in a server update
 - eg, Updating sort code and account number results in 2 updates on server
- Makes things more difficult to debug
- Causes more network traffic
- Makes it more difficult to handle faults and conflicts
 - What if I update two properties, one is successful, other fails – have to manage rollbacks?
- Update the properties on your [Managed] Objects in a single, atomic transaction
- Our preference is to set autoCommit to 'false' on the ServiceLocator

```
<mx:DataService  
  id="projectService"  
  destination="projectService"  
  autoCommit="false" />
```

- When you are ready, call commit() on the data service

- Auto-commit false results in changes being batched up
 - Kept in client-side data store
- You must call commit() for:
 - createItem()
 - deleteItem()
 - Updates to read-write properties
- You don't need to call commit() for:
 - fill()
 - count()
- Need to track changes?
 - createItem() and deleteItem() return an AsyncToken to track item-by-item changes.
 - A commit() returns an AsyncToken which is a "master" token for the batch as a whole.

- Always add a fault handler to data services
- Unexpected faults can happen and should be handled
 - Makes debugging much quicker
- Consider how it might affect the user experience
 - How do you inform user that something has gone wrong?

Data Management Services Hints and Tips



Which properties on a class are managed?



- [Managed] can only be used on class definitions
 - Cannot be used on properties
- Any modification to a read-write properties on the class will result in a server update
- If you mark a read-write property as [Transient] it will not be managed
- Lets you add client-side only properties to your managed objects

```
package com.adobe;
{
    [Managed]
    [RemoteClass(alias="AccountVO")]
    public class Account
    {
        public var sortCode : Number;
        public var accountNumber : String;

        [Transient]
        public var hashCode : String;
    }
}
```

Why does deleteItem() result in other updates?



- If you call deleteItem()...
 - You may see multiple updates on the server
- The framework checks to see if the item is referenced elsewhere
- If so, updates are issued to remove the item from those properties
 - This is done before item is deleted

How do conflicts apply to associations?



- Associations can be set up in two main ways
 - Hierarchical values approach (single data management service destination)
 - Managed association approach (define parent destination and associations to child destinations)
- Consider this “hierarchical values” example, which defines one managed service

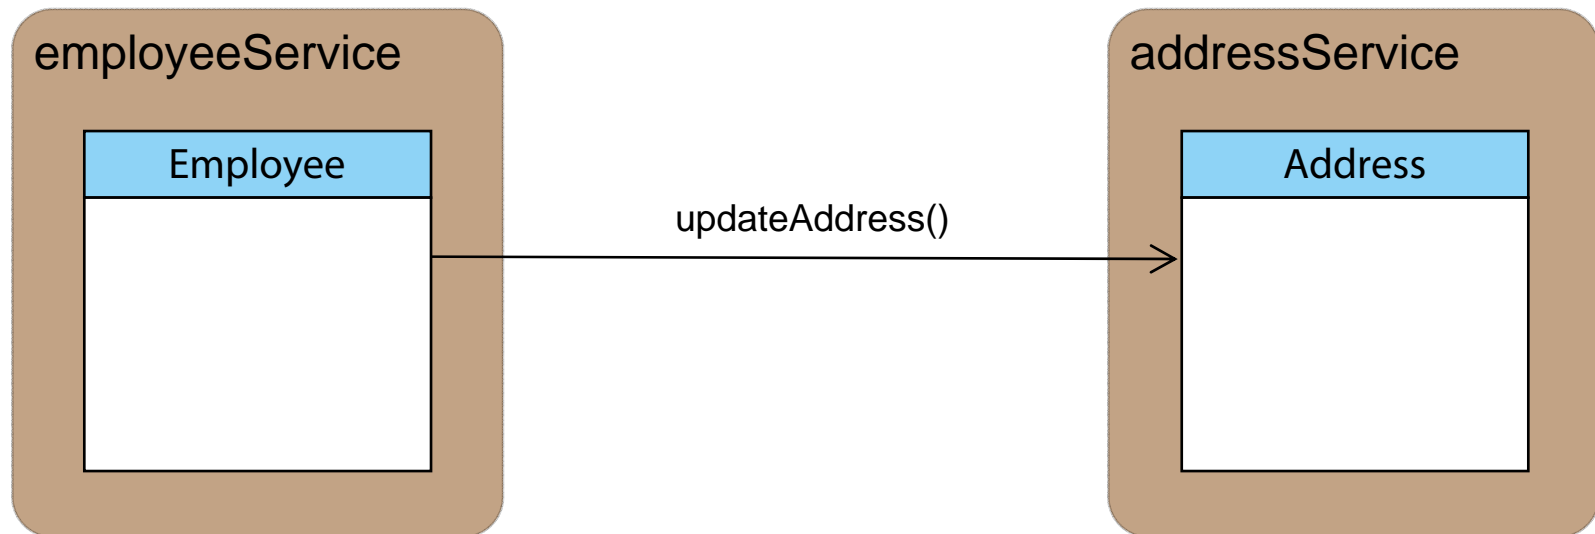


- The parent assembler (`employeeService`) will handle conflict from both classes

How do conflicts apply to associations?



- If you are using an managed association approach
 - ie. Define parent destination and associations to manage a tree of nested objects
- By default, each service handles its own conflicts
- But if you don't create a data service, an internal one is created for you
 - But you can't then handle any conflicts
- You should explicitly create data services and handle conflicts



- Look at the code below
 - It is a [Managed] class with a [Bindable] getter/setter
 - Setter is responsible for dispatching its own change event
- You may expect this to work...
 - Actually causes a stack overflow

```
[Managed]
[RemoteClass(alias="com.adobe.Folder")]
public class Folder
{
    [Bindable(event="foldersChange")]
    public function get folders() : ListCollectionView
    {
        ...
    }

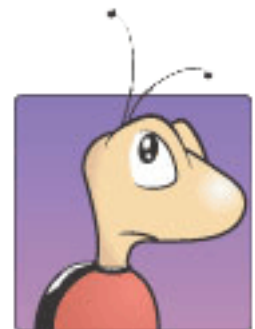
    public function set folders( value : ListCollectionView ) : void
    {
        ...
        dispatchEvent( new Event(" foldersChange" ) );
    }
}
```

Why do I get a stack overflow if I use [Bindable]?



- When you specify [Bindable(event="...")] for a property on a class which has a class level [Bindable] tag, it indicates that this property is handling its own change event
 - The compiler knows not to codegen the property change event from the class level [Bindable] metadata
- The [Managed] tag implicitly makes a class [Bindable] too
- But, the compiler does not recognise this so it will still do the codegen
 - It causes the stack overflow error
 - This is a bug
- The workaround is to remove the [Managed] tag and do the coding for the managed support yourself.
 - http://help.adobe.com/en_US/lifecycle/es/lcds_dev_guide.pdf (page 240-241)
 - Alternatively, keep the generated source from the compiler

Bug#: 169885
(fixed)



Why is the setter called when I call the getter?



- You may see times where your setter is called when you don't expect it
 - Noticeable when debugging data management services
- Underlying getter is calling the setter
- Managed object auto-generated code does this
- Doesn't cause any problems, but be aware of it

```
[Bindable(event="propertyChange")]  
public function get firstName():String  
{  
    _132835675firstName = mx.data.utils.Managed.getProperty(this, "firstName", _132835675firstName);  
    return _132835675firstName;  
}
```

How do you keep data on the client after logout?



- When a user logs out, their session is invalidated (for security reasons)
- Causes a `release()` call on `DataService`
 - Results in calls to `releaseCollection(coll, true)` for all managed collections
 - This releases and clears data in the client side collections too
- May want to keep some data on the client after logout (eg, localisation data)
- To keep the data on the client after logout, use the following approach:

// this prevents the server from storing references to any subsequent fills on this service

```
localeDataService.autoSyncEnabled = false;  
localeDataService.fill( localeCollection );
```

// wait for the result of the fill...

```
localeDataService.releaseCollection( localeCollection, false );
```

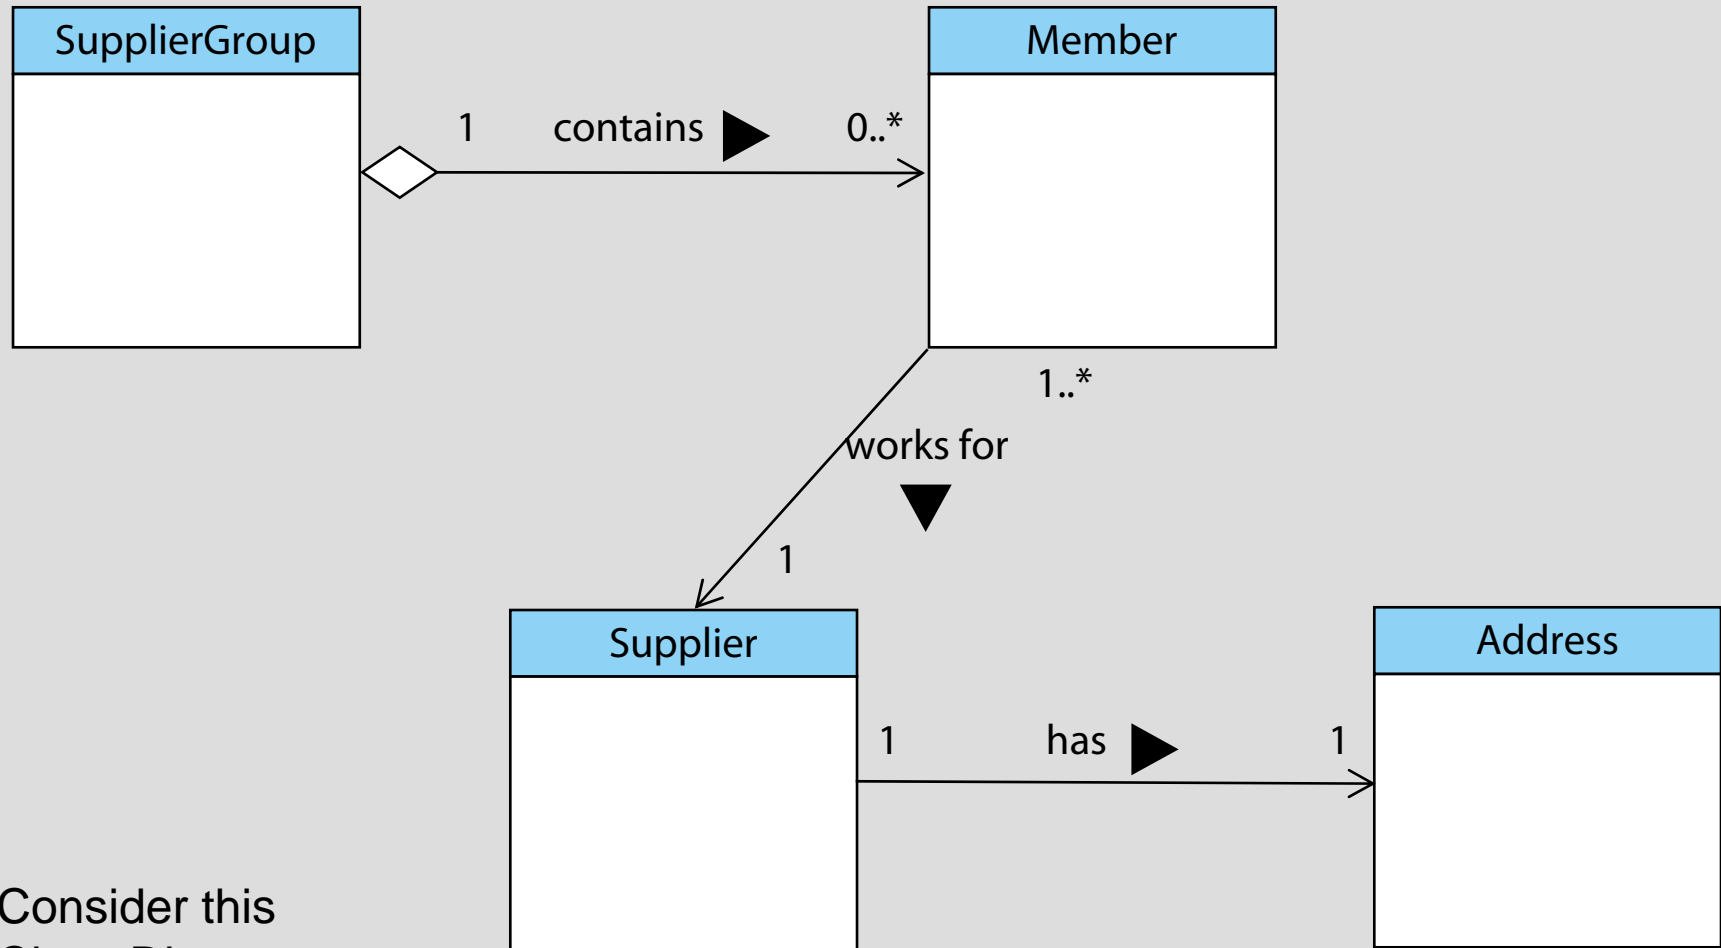
- The second parameter to `releaseCollection()` indicates whether to clear the collection
- Now, on logout, the collection won't be released, as it has already been released
- Note, however, this approach makes a copy of the objects
 - Any subsequent updates will not be pushed to the client
 - Effectively, they are no longer managed objects

If you sort a column on a DataGrid what happens?



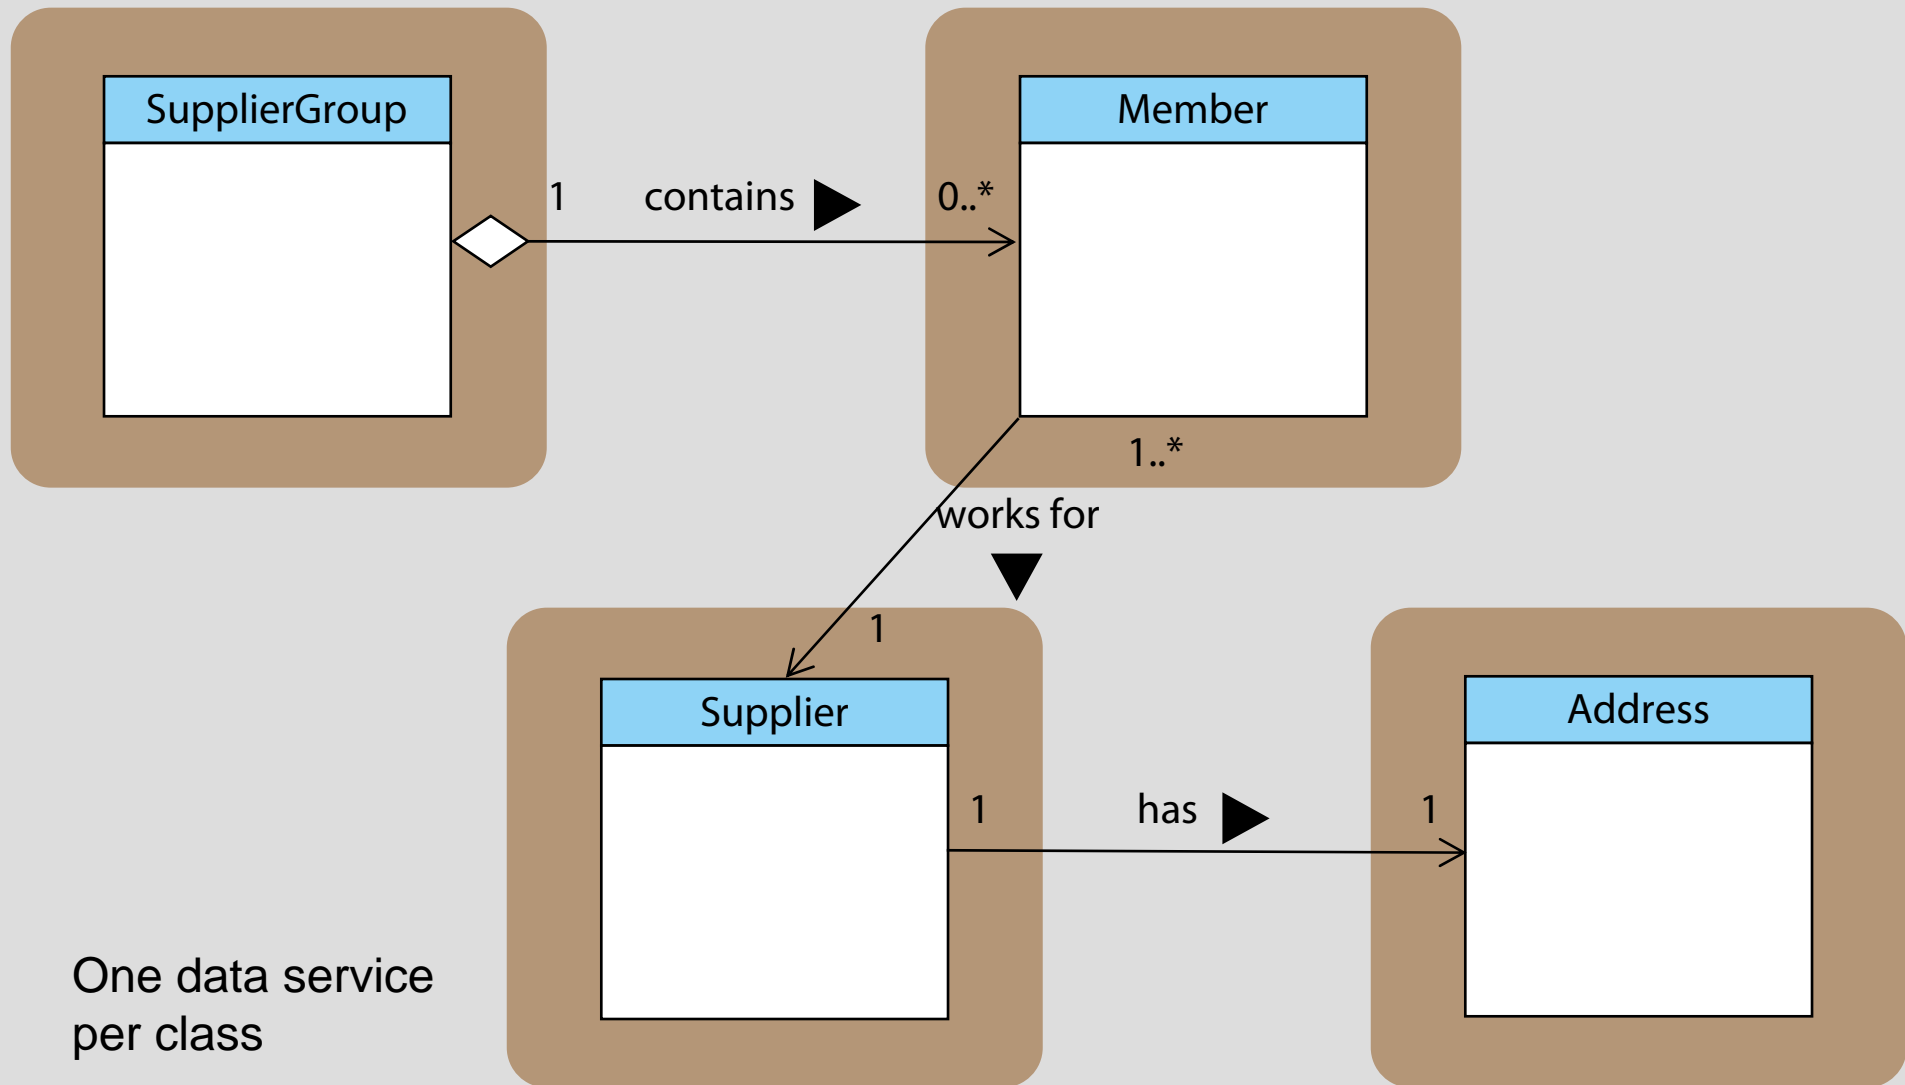
- Imagine we have a DataGrid with a paged remote collection
- If the user chooses to sort a column, all the data will be retrieved before sorting
 - Sorting is done client-side
- Solution is to write a custom handler for the headerRelease event
- You then call fill() on the DataService
 - Pass the sort criteria as a parameters to the fill

How clever is Data Management Services?



Consider this
Class Diagram

Which maps on to the following data services...



How do I add new member?



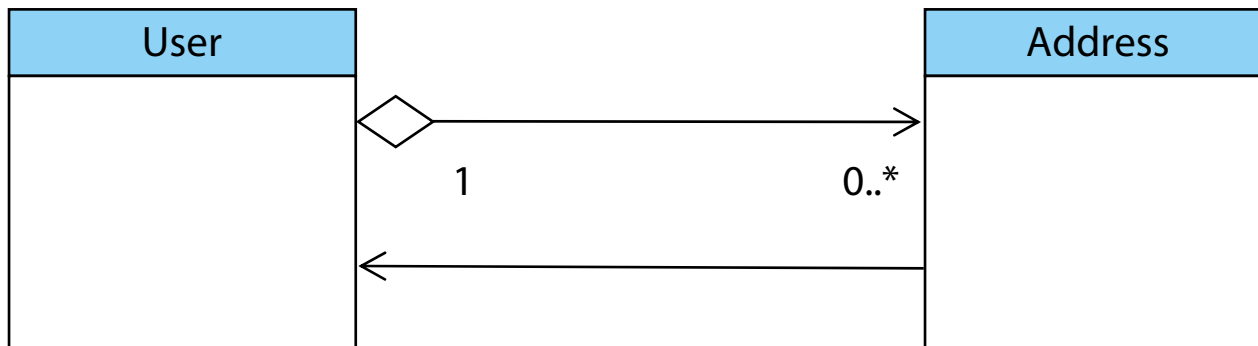
- Use Case
 - We want to add a new member to a supplier group
 - The member works for a new supplier, who is not yet known to the system
 - Which means we need to add a new address too
- Just how clever is Data Management Services?
 - Can I simply create the new member, with a new supplier and a new address and add it to the collection of members in the supplier group?
 - Or do I have to explicitly create the items through individual services?
 - If so, do I need to chain the create items together, creating the address first, then the supplier and finally the member?
 - ie. Do I have to start from the item with no dependencies (the address)

- Data management services will create items in the right order
- Calls `createItem()` on the object chain
 - Starts with the one with no dependencies
 - Creates address, then supplier, then member then updates the supplier group
- So, Data Management Services is really clever!
 - But needs the appropriate associations defined in configuration files
- Only one object in the chain needs to be managed
 - If others are not managed, data services will be created for them automatically
 - But then, we can't handle conflicts or faults
- What if you have a bi-directional relationship?
 - In this case, data management services cannot determine the dependencies
 - It needs a clue, so we set `read-only="true"` on the item with no dependency
 - Example coming...

Consider this class diagram...



◀ `<many-to-one property="user" destination="userService" lazy="true" />`



`<one-to-many property="addresses" destination="addressService" read-only="true" lazy="true" />` ▶

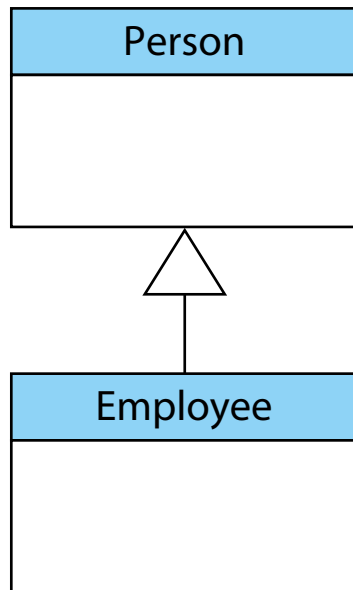
- User and Address have a bi-directional relationship
- Specify read-only="true" on addressService
- Indicates that Address has no dependency on User
- Address will be added first

*A 10 point footnote can go here, if necessary

Why do I get a type coercion error?



- Always ensure your managed classes are linked into the SWF
 - If you don't an object of type "ManagedObjectProxy" will be created.
- In this example, we call a service that returns an Employee object
- If Employee isn't linked into the SWF and we try to coerce to Person, we get a coercion error



```
selectedItem = mx.data.ManagedObjectProxy (@175bc711)
  address = Address (@1751a479)
  data = undefined
  department = "Adobe Consulting"
  destination = "employeeService"
  firstName = "Xavier"
  label = undefined
  lastName = "Agnetti"
  object = Object (@174cc5b1)
  personId = 4
  referencedIds = Object (@174df1
  type = null
  uid = "employeeService:::4"
```

```
selectedItem = Employee (@17421ec1)
  address = Address (@175098b1)
  department = "Adobe Consulting"
  destination = "employeeService"
  firstName = "Tom"
  lastName = "Sugden"
  personId = 2
  referencedIds = Object (@174b3a39)
  uid = "employeeService:::2"
```

What is the “highlander principle”?

- Taken from the film “Highlander”, which was released in 1986.
- “There can be only one” - there can be only one instance of an object with each id
- Data management services will ensure there is only one copy of an object on the client.
- But this does depend on how you setup your associations
 - Best practice is to use managed associations
 - Hierarchical associations do not assure the highlander principle



- MFG.com project
- Cairngorm Enterprise
 - Why and what
- Data Management Services
 - Best practices
 - Hints and tips
- Next for Cairngorm Enterprise?
 - We're looking to open up discussions to the community
 - There is a cairngorm-devel mailing list on Yahoo Groups
 - Cairngorm Committee?

Any questions?



Questions?

“There can be only one”

MAX

ADOBE® CONSULTING

